

Quiz: Indexes (2)

Collaboration policy:

You may NOT:

1. discuss the exam with any human other than Mike; this includes:
 - (a) asking your friend for clarification about what a problem is asking
 - (b) asking your friend if they've finished working on the quiz
 - (c) posting questions to github, stackoverflow, or any other online forum

You may:

1. take as much time as needed
2. use any written notes / electronic resources you would like (including postgres documentation and AI systems like ChatGPT)
3. use the lambda server and connect to a running postgres instance
4. ask Mike to clarify questions via email

Name:

All of the problems assume the following schema.

```
CREATE TABLE users (  
    id_users BIGINT PRIMARY KEY,  
    created_at TIMESTAMPTZ,  
    username TEXT NOT NULL,  
    description TEXT  
);  
  
CREATE TABLE tweets (  
    id_tweets BIGINT UNIQUE NOT NULL,  
    id_users BIGINT REFERENCES users(id_users),  
    in_reply_to_user_id BIGINT REFERENCES users(id_users),  
    created_at TIMESTAMPTZ,  
    text TEXT  
);  
  
CREATE TABLE tweet_tags (  
    id_tweets BIGINT REFERENCES tweets(id_tweets),  
    tag TEXT,  
    PRIMARY KEY(id_tweets, tag)  
);
```

1. Recall that certain constraints create indexes on the appropriate columns. List the equivalent `CREATE INDEX` commands that are run by these constraints.

Note: When solving all subsequent problems, assume only the indexes above have been created. For example, if problem 3 asks you to create an index, you should not assume that index has been created when working on problem 4.

2. Create index(es) so that the following query will run as efficiently as possible.

Do not create any unneeded indexes; if no new indexes are needed, say so.

```
SELECT count(*)
FROM users
WHERE (id_users >= :min_id_users)
      OR (username >= :min_username AND username < :max_username)
```

3. Create index(es) so that the following query will run as efficiently as possible.

Do not create any unneeded indexes; if no new indexes are needed, say so.

```
SELECT count(*)
FROM users
WHERE username=lower(:username);
```

4. The following SQL query returns tweets where either the text or the description of the user match a FTS query.

```
SELECT id_tweets
FROM tweets, users
WHERE ( to_tsvector('english', text)
      || to_tsvector('english', description)
      )
      @@ to_tsquery('english', :query)
;
```

The query cannot be sped up using indexes because the left hand side of the @@ operator references columns from two separate tables, and we cannot create an index on two separate tables.

Rewrite the query above into an equivalent query that can be sped up with indexes, and also provide any needed indexes.